

```

import tensorflow as tf
tf.test.gpu_device_name()

'/device:GPU:0'

import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from tensorflow.keras.utils import to_categorical
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense, Flatten
from tensorflow.keras.optimizers import Adam
from sklearn.impute import SimpleImputer

# Load the datasets
train_data = pd.read_csv('mnist_train.csv')
test_data = pd.read_csv('mnist_test.csv')

# Explore the dataset
print("Number of unique classes:", train_data['label'].nunique())
print("Number of features:", len(train_data.columns) - 1) # Excluding the label column
print("Missing values in training set:", train_data.isnull().sum().sum())

#X_test->test_images, y_test->test_labels

# Handle missing values by replacing NaN with mean
X_train = train_data.iloc[:, 1:].values
X_test = test_data.values[:, 1:]

# Use SimpleImputer to replace NaN with mean
imputer = SimpleImputer(strategy='mean')
X_train = imputer.fit_transform(X_train)
X_test = imputer.transform(X_test)

# Convert to float32 explicitly
X_train = X_train.astype('float32')
X_test = X_test.astype('float32')

# Normalize pixel values
X_train /= 255.0
X_test /= 255.0

# Reshape images
X_train = X_train.reshape(-1, 28, 28, 1)
X_test = X_test.reshape(-1, 28, 28, 1)

# Encode labels
y_train = to_categorical(train_data['label'])
y_test = to_categorical(test_data['label'])


# Split into training and validation sets
X_train, X_test, y_train, y_test = train_test_split(
    X_train, y_train, test_size=0.2, random_state=42
)

# Visualize some images
fig, axes = plt.subplots(2, 5, figsize=(10, 5))
for i, ax in enumerate(axes.flat):
    ax.imshow(X_train[i].reshape(28, 28), cmap='gray')
    ax.set_title(f"Label: {train_data['label'][i]}")
    ax.axis('off')
plt.show()

```



Number of unique classes: 10  
 Number of features: 784  
 Missing values in training set: 0



```
from sklearn.metrics import confusion_matrix
```

```
#Code steps
```

```
#Builds a neural network model.
#Compiles the model.
#Trains the model on training data.
#Evaluates the model on the test set.
#Prints the test accuracy.
#Calculates and prints the confusion matrix.
```

```
# Experiment 1
```

```
# Build the ANN model
model_1 = Sequential()
model_1.add(Flatten(input_shape=(28, 28, 1))) # Flatten the 28x28 images
model_1.add(Dense(128, activation='relu'))
model_1.add(Dense(10, activation='softmax')) # Output layer with 10 classes
#The output layer is added with 10 neurons (representing the 10 classes in the MNIST dataset)
```

```
# Compile the model
model_1.compile(optimizer=Adam(learning_rate=0.001), loss='categorical_crossentropy', metrics=['accuracy'])
```

```
# Train the model
Train_model_1 = model_1.fit(X_train, y_train, epochs=10, batch_size=64, validation_data=(X_test, y_test))
```

```
# Evaluate on test set
test_loss_1, test_accurecy_1 = model_1.evaluate(X_test, y_test)
print(f"\nTest Accuracy: {test_accurecy_1}\n")
```

```
# Confusion Matrix
y_predict1 = model_1.predict(X_test)
Confusion_matrix_1 = confusion_matrix(np.argmax(y_test, axis=1), np.argmax(y_predict1, axis=1))
print("Confusion Matrix (Experiment 1):\n", Confusion_matrix_1)
```

```
Epoch 1/10
750/750 [=====] - 4s 5ms/step - loss: 0.3308 - accuracy: 0.9076 - val_loss: 0.1904 - val_accuracy: 0.9473
Epoch 2/10
750/750 [=====] - 3s 4ms/step - loss: 0.1569 - accuracy: 0.9545 - val_loss: 0.1369 - val_accuracy: 0.9617
Epoch 3/10
750/750 [=====] - 3s 4ms/step - loss: 0.1096 - accuracy: 0.9674 - val_loss: 0.1124 - val_accuracy: 0.9682
Epoch 4/10
750/750 [=====] - 2s 3ms/step - loss: 0.0835 - accuracy: 0.9764 - val_loss: 0.0973 - val_accuracy: 0.9721
Epoch 5/10
750/750 [=====] - 3s 4ms/step - loss: 0.0648 - accuracy: 0.9809 - val_loss: 0.0942 - val_accuracy: 0.9713
Epoch 6/10
750/750 [=====] - 3s 4ms/step - loss: 0.0528 - accuracy: 0.9846 - val_loss: 0.0924 - val_accuracy: 0.9714
Epoch 7/10
750/750 [=====] - 3s 4ms/step - loss: 0.0423 - accuracy: 0.9873 - val_loss: 0.0843 - val_accuracy: 0.9758
Epoch 8/10
750/750 [=====] - 3s 4ms/step - loss: 0.0344 - accuracy: 0.9899 - val_loss: 0.0860 - val_accuracy: 0.9748
Epoch 9/10
750/750 [=====] - 2s 3ms/step - loss: 0.0274 - accuracy: 0.9925 - val_loss: 0.0900 - val_accuracy: 0.9740
Epoch 10/10
750/750 [=====] - 3s 4ms/step - loss: 0.0231 - accuracy: 0.9935 - val_loss: 0.0849 - val_accuracy: 0.9756
375/375 [=====] - 1s 2ms/step - loss: 0.0849 - accuracy: 0.9756
```

```
Test Accuracy: 0.9755833148956299
```

```
375/375 [=====] - 1s 2ms/step
Confusion Matrix (Experiment 1):
[[1156   0    3    1    4    1    6    0    3    1]
 [   0 1311    3    2    2    0    1    1    1    1]
 [   1    8 1145    6    2    0    1    6    3    2]
 [   0    1   14 1179    0   11    2    4    6    2]
```

```
[ 1  3  0  1 1154  1  3  1  1 11]
[ 7  3  2 11  4 1058  9  2  6  2]
[ 1  0  3  0  1  1 1171  0  0  0]
[ 1  7  9  1  3  0  0 1275  1  2]
[ 3  7  3  6  5  6  8  1 1114  7]
[ 4  1  1  3 19  2  1 12  7 1144]]
```

#Hyperparameters:

```
#Learning Rate: Determines the step size during optimization.
#Batch Size: Number of training examples used in one iteration.
#Number of Layers and Neurons: Architecture of the neural network.
#Activation Functions: Choice of activation functions in each layer.
#Dropout Rate: Regularization technique to prevent overfitting.
#Weight Initialization: Initial values assigned to the weights.
```

# Experiment 2

# Build the ANN model

```
model_2 = Sequential()
model_2.add(Flatten(input_shape=(28, 28, 1)))
model_2.add(Dense(256, activation='relu')) # Different number of neurons
model_2.add(Dense(10, activation='softmax'))
```

# Compile the model

```
model_2.compile(optimizer=Adam(learning_rate=0.01), loss='categorical_crossentropy', metrics=['accuracy'])
```

# Train the model

```
Train_model_2 = model_2.fit(X_train, y_train, epochs=10, batch_size=32, validation_data=(X_test, y_test))
```

# Evaluate on test set

```
test_loss_2, test_accurecy_2 = model_2.evaluate(X_test, y_test)
print(f"\nTest Accuracy: {test_accurecy_2}\n")
```

# Confusion Matrix

```
y_predict2 = model_2.predict(X_test)
Confusion_matrix_2 = confusion_matrix(np.argmax(y_test, axis=1), np.argmax(y_predict2, axis=1))
print("Confusion Matrix (Experiment 2):\n", Confusion_matrix_2)
```

Epoch 1/10

1500/1500 [=====] - 5s 3ms/step - loss: 0.2602 - accuracy: 0.9237 - val\_loss: 0.1891 - val\_accuracy: 0.9444

Epoch 2/10

1500/1500 [=====] - 6s 4ms/step - loss: 0.1685 - accuracy: 0.9540 - val\_loss: 0.1899 - val\_accuracy: 0.9513

Epoch 3/10

1500/1500 [=====] - 5s 3ms/step - loss: 0.1500 - accuracy: 0.9598 - val\_loss: 0.1716 - val\_accuracy: 0.9598

Epoch 4/10

1500/1500 [=====] - 5s 3ms/step - loss: 0.1367 - accuracy: 0.9641 - val\_loss: 0.2043 - val\_accuracy: 0.9533

Epoch 5/10

1500/1500 [=====] - 6s 4ms/step - loss: 0.1172 - accuracy: 0.9697 - val\_loss: 0.2140 - val\_accuracy: 0.9544

Epoch 6/10

1500/1500 [=====] - 5s 3ms/step - loss: 0.1168 - accuracy: 0.9707 - val\_loss: 0.2829 - val\_accuracy: 0.9479

Epoch 7/10

1500/1500 [=====] - 6s 4ms/step - loss: 0.1114 - accuracy: 0.9724 - val\_loss: 0.2571 - val\_accuracy: 0.9526

Epoch 8/10

1500/1500 [=====] - 5s 3ms/step - loss: 0.1044 - accuracy: 0.9760 - val\_loss: 0.2380 - val\_accuracy: 0.9628

Epoch 9/10

1500/1500 [=====] - 5s 3ms/step - loss: 0.1070 - accuracy: 0.9762 - val\_loss: 0.2119 - val\_accuracy: 0.9664

Epoch 10/10

1500/1500 [=====] - 6s 4ms/step - loss: 0.0877 - accuracy: 0.9795 - val\_loss: 0.2582 - val\_accuracy: 0.9609

375/375 [=====] - 1s 3ms/step - loss: 0.2582 - accuracy: 0.9609

Test Accuracy: 0.9609166383743286

375/375 [=====] - 1s 2ms/step

Confusion Matrix (Experiment 2):

```
[[1151  0  1  0  2  2  5  0 13  1]
[  0 1294 11  2  1  0  1  3  9  1]
[  2  3 1136  8  0  1  2 13  8  1]
[  0  0  10 1173  0  7  0  8 20  1]
[  3  1  2  0 1106  1 23 10 12 18]
[  7  2  2 32  1 1012 18  2 26  2]
[  3  0  2  0  1  1 1163  1  6  0]
[  1  4 10  3  1  0  0 1273  5  2]
[  5  1  7  5  2  2  5  1 1132  0]
[  5  1  2 24 21  2  1 26 21 1091]]
```

```

from sklearn.model_selection import GridSearchCV
from sklearn.neighbors import KNeighborsClassifier
from sklearn.metrics import accuracy_score, classification_report, confusion_matrix

# Flatten images
X_train_flat = X_train.reshape(X_train.shape[0], -1)
X_test_flat = X_test.reshape(X_test.shape[0], -1)

# Convert one-hot encoded labels back to integers
y_train_int = np.argmax(y_train, axis=1)
y_test_int = np.argmax(y_test, axis=1)

# Define the K-NN model
knn_model = KNeighborsClassifier()

# Define hyperparameter grid for grid search
param_grid = {'n_neighbors': [3, 5, 7], 'weights': ['uniform', 'distance']}

# Perform grid search
grid_search = GridSearchCV(knn_model, param_grid, cv=3, scoring='accuracy')
grid_search.fit(X_train_flat, y_train_int)

# Print the best parameters
print("Best Parameters:", grid_search.best_params_)

# Train the model with the best parameters
best_knn_model = grid_search.best_estimator_
best_knn_model.fit(X_train_flat, y_train_int)

# Predict on the validation set
knn_val_predictions = best_knn_model.predict(X_test_flat)

# Evaluate the model on the validation set
test_accuracy_3 = accuracy_score(y_test_int, knn_val_predictions)
print(f"Validation Accuracy (K-NN): {test_accuracy_3}")

# Print classification report and confusion matrix
print("\nClassification Report:")
print(classification_report(y_test_int, knn_val_predictions))

print("\nConfusion Matrix:")
print(confusion_matrix(y_test_int, knn_val_predictions))

```

```

Best Parameters: {'n_neighbors': 3, 'weights': 'distance'}
Validation Accuracy (K-NN): 0.9735833333333334

```

```

Classification Report:

```

	precision	recall	f1-score	support
0	0.98	0.99	0.99	1175
1	0.96	1.00	0.98	1322
2	0.99	0.96	0.97	1174
3	0.97	0.97	0.97	1219
4	0.97	0.97	0.97	1176
5	0.97	0.97	0.97	1104
6	0.98	0.99	0.99	1177
7	0.97	0.98	0.97	1299
8	0.98	0.95	0.97	1160
9	0.96	0.96	0.96	1194
accuracy			0.97	12000
macro avg	0.97	0.97	0.97	12000
weighted avg	0.97	0.97	0.97	12000

```

Confusion Matrix:
[[1167  0  0  0  1  3  1  1  2]
 [  0 1316  1  1  1  0  0  2  0  1]
 [  7  11 1131  2  3  1  0 15  3  1]
 [  1  0  8 1181  0 13  0  4  6  6]
 [  0  6  1  1 1135  0  2  4  0 27]
 [  5  3  0 14  2 1067  8  0  4  1]
 [  3  2  0  0  1  3 1168  0  0  0]
 [  0 19  1  1  2  0  0 1268  3  5]
 [  3  8  4 15  6 12  3  3 1100  6]
 [  4  2  2  1 15  2  2 15  1 1150]]

```

```
!pip install joblib
```

Requirement already satisfied: joblib in /usr/local/lib/python3.10/dist-packages (1.3.2)

```

from tensorflow.keras.models import load_model

import joblib
from sklearn.metrics import confusion_matrix

# Compare the outcomes
if test_accurecy_1 > test_accurecy_2 and test_accurecy_1 > test_accurecy_3:
    BestModel = model_1
    print("Experiment 1 has the higher accuracy on the validation set. ANN model_1\n")
elif test_accurecy_2 > test_accurecy_1 and test_accurecy_2 > test_accurecy_3:
    BestModel = model_2
    print("Experiment 2 has the higher accuracy on the validation set. ANN model_2\n")
else:
    BestModel = knn_model
    print("Experiment 3 has the higher accuracy on the validation set. KNN model\n")

# Save the best model
if isinstance(BestModel, KNeighborsClassifier):
    joblib.dump(BestModel, "Best_model_knn.joblib")
else:
    BestModel.save("Best_model_Ann.h5")

# Reload the model
if isinstance(BestModel, KNeighborsClassifier):
    loaded_model = joblib.load("Best_model_knn.joblib")
else:
    loaded_model = load_model("Best_model_Ann.h5")

# Display the model summary
if not isinstance(BestModel, KNeighborsClassifier):
    loaded_model.summary()

# Evaluate the reloaded model on the test set
if isinstance(BestModel, KNeighborsClassifier):
    knn_val_predictions = loaded_model.predict(X_test)
    accuracy = accuracy_score(y_test, knn_val_predictions)
    print(f"\nTest Accuracy using the reloaded KNN model: {accuracy}")
    # Get confusion matrix
    cm = confusion_matrix(y_test, knn_val_predictions)
    print("\nConfusion Matrix:")
    print(cm)
else:
    test_loss, test_accuracy = loaded_model.evaluate(X_test, y_test)
    print(f"\nTest Accuracy using the reloaded ANN model: {test_accuracy}")

Experiment 1 has the higher accuracy on the validation set. ANN model_1

Model: "sequential_6"

Layer (type)                 Output Shape                 Param #
=====
flatten_6 (Flatten)          (None, 784)                  0

dense_12 (Dense)              (None, 128)                  100480

dense_13 (Dense)              (None, 10)                   1290

=====
Total params: 101770 (397.54 KB)
Trainable params: 101770 (397.54 KB)
Non-trainable params: 0 (0.00 Byte)

375/375 [=====] - 2s 5ms/step - loss: 0.0849 - accuracy: 0.9756

Test Accuracy using the reloaded ANN model: 0.9755833148956299

```

